

Load data:

```
10 # import data
11 (X_train, y_train), (X_test, y_test) = mnist.load_data()
12
13 # scale data
14 X_train = X_train / 255
15 X_test = X_test / 255
```

Generate noisy data:

```
17 # generate noisy data
18 X_train_noisy = np.clip(X_train
19     + np.random.normal(loc=0.0, scale=0.5, size=X_train.shape), 0., 1.)
20 X_test_noisy = np.clip(X_test
21     + np.random.normal(loc=0.0, scale=0.5, size=X_test.shape), 0., 1.)
```

Build convolutional autoencoder--undercomplete to extract the most useful features. The autoencoder reconstructs noisy images as clean images:

```
8 # build convolutional autoencoder
9 conv_encoder = models.Sequential([
10     layers.Reshape([28, 28, 1], input_shape=[28, 28]),
11     layers.Conv2D(16, (3, 3), padding='same', activation='selu'),
12     layers.MaxPool2D((2, 2)),
13     layers.Conv2D(8, (3, 3), padding='same', activation='selu'),
14     layers.MaxPool2D((2, 2))
15 ])
16
17 conv_decoder = models.Sequential([
18     layers.Conv2D(8, (3, 3), padding='same', activation='selu',
19         input_shape=[7, 7, 8]),
20     layers.UpSampling2D((2, 2)),
21     layers.Conv2D(4, (3, 3), padding='same', activation='selu'),
22     layers.UpSampling2D((2, 2)),
23     layers.Conv2D(1, (3, 3), padding='same', activation='sigmoid'),
24     layers.Reshape([28, 28])
25 ])
26
27 conv_ae = models.Sequential([conv_encoder, conv_decoder])
28 conv_ae.compile(loss='binary_crossentropy', optimizer='adam')
29 history = conv_ae.fit(X_train_noisy, X_train, epochs=20,
30     validation_data=(X_test_noisy, X_test))
```

Results--noisy images (top) and clean autoencoder reconstructions (bottom):

